ACCOUNT DEBRIEF · AI OBSERVABILITY

Acme: winning back a former customer through AI governance

Discovery analysis, account history, market intelligence and the demo + POC path. Prepared for Dynatrace leadership on April 8, 2026.

DYNATRACE INTERNAL · CONFIDENTIAL | DEMO: MON APR 27, 10–11 AM PT

THE BOTTOM LINE

Acme was a Dynatrace customer from 2015 to 2017 (about $42K ACV). It left over **pricing** and **AWS coverage**, and runs Datadog today. **Marcus Hale**, Acme's Director of AI, has to **show executives ROI on AI** and **govern a fast-growing agent population** with a tiny team. He has **written our demo agenda for us** (6 asks). Win Apr 27 by answering those asks with live AI Observability tied to business outcomes. Then land a **30-day AI Value & Governance POV** on one real AI workload.

3 : 1
Humans to AI agents today. Hale expects it to **flip in 2027**.

6
Explicit demo asks Hale put in writing on Apr 8

$42K
Prior ACV (2015–17, 24 host units). Lost on price and AWS gaps.

1
Person at Acme with "AI" in the title. Single-threaded today.

Prove AI value

Executives went from "use it as much as you can" to *"you spent $8,000 in credits, so what did we get?"* He needs ROI, cost and **agent productivity**, not just agent health.

Govern autonomous AI

Identity & access for agents "just as much as for humans", guardrails, a **kill switch**, and a QA gate for vibe-coded software before production.

Run AI reliably at scale

End-to-end visibility into agents and workflows that span **multiple systems and models**, from experiment to "what good looks like" in a real enterprise.

Recommended moves

1. Run the Apr 27 demo **on his six asks, in his order**. No platform tour.
2. **Fix three over-promises before the demo** (kill switch, agent IAM, the "your models" slide).
3. Offer a **30-day POV** with five measurable success criteria and an exec readout that **Hale presents**.
4. **Multi-thread now**: Digital VP, Event Technology, InfoSec, IT Engineering. Bring a Dynatrace exec sponsor.

60-SECOND READ

Executive summary

- **The opportunity:** re-enter a marquee former customer through **net-new AI scope**, not a rip-and-replace. Salesforce's P2B model estimates **about $310K ARR potential** for Acme, versus the **$42K ACV we lost in 2017**.
- **The buyer:** Marcus Hale, Acme's Director of AI, is under **executive pressure to prove AI ROI** and govern an agent population he expects to outnumber employees, probably in 2027.
- **The ask:** he wrote our agenda. There are **six asks**, led by "what good looks like", plus cost, reliability, multi-model visibility, governance and business outcomes.
- **The play:** win the **Apr 27 demo** → a **30-day AI Value & Governance proof of value** → expand into race-day reliability and logs → platform win-back from Datadog (via a reseller).
- **The main threat: Datadog LLM Observability**, switched on inside the incumbent contract with no new vendor review. We win on causal root cause across agents, business-outcome linkage and AI-eval partner evaluations (acquisition now closed).
- **The open questions:** whether a custom AI app exists to instrument (the confirmed estate is an enterprise chat assistant), who funds it, and deal size. The opp shows $0 ACV and a Dec 2026 close today.

~$310K
ARR potential (Salesforce P2B estimate)

Apr 27
Demo, 10–11 AM PT. Agenda written by the buyer.

30 days
Proposed proof of value with 5 success criteria

What we need from leadership

Executive sponsor
A Dynatrace executive paired with a Acme brand/digital executive after Apr 27

Principal SE coverage
An AI-observability principal ready to run the proof of value

Commercial flexibility
An AI-scoped subscription land with no caps. We lost twice on price in 2017.

Fix before Apr 27

- Verify the **kill switch** and **the partner AI gateway** capabilities with product before claiming them
- On **agent identity**, say Dynatrace **observes and audits**. Enforcement belongs to IAM partners.
- Correct the deck: the model list (Gemini, OpenAI, Microsoft, Claude) is **Dynatrace's**, not Acme's
- Tell Acme the AI-eval acquisition has **closed**. On the call we said it was pending.

IN PRIORITY ORDER · BUSINESS OUTCOMES FIRST

Main talking points

Each point leads with what Acme gets. The capability comes second.

#	TALKING POINT	BUSINESS OUTCOME FOR ACME	POTENTIAL FOR DYNATRACE
1	**"Here's what good looks like."** A target end state for AI in a real enterprise, with Acme placed on the maturity ladder.	A **clear roadmap from experiments to operating AI at scale**. This is his stated top angle.	Frames us as the AI operating partner, not a tool
2	**"Prove what your AI spend buys."** Cost per workflow, tied to one outcome KPI on an executive scorecard.	Answers "eight grand in credits… so what did we get for it?" with **defensible AI ROI**	Core land: AI Observability
3	**"Govern agents like employees."** For every sanctioned agent: actions, data access and cost, all auditable, with PII masked at the edge.	**Scale autonomy without risking fan trust**. Gives his "AI intake lead" gate evidence after approval.	Governance anchor; AI-eval partner evaluations; pipeline attach
4	**"Run AI at scale with a scrappy team."** Automatic root cause across multi-model, multi-system agent workflows.	**Scale output through AI** without adding headcount. Agent failures resolved in minutes, not hours.	Stickiness: AI becomes an ops dependency
5	**"From experiment to scale, proven."** Customer references plus Forrester TEI (466% ROI, payback under 6 months).	**Lower decision risk**, and a story Hale can take to his executives and on stage	Faster land; reference potential
6	**"Ship AI-built software safely."** Production-aware guardrails for vibe-coded changes (once verified).	**Speed without incidents** as AI-generated code reaches production	AI SRE / developer expansion
7	**"Protect race day and the fan experience."** One platform covers race data (edge Kubernetes → AWS), apps and streaming.	**Competition integrity and fan experience** on the sport's highest-stakes days	Expansion: Racing Ops / Event Tech
8	**"Control observability cost as AI multiplies telemetry."** Pipeline-level volume control, then consolidation.	**Predictable cost** as AI workloads grow data volumes	Datadog displacement via the reseller opp

Say

- "Manage agents like staff. Here's the evidence for that."
- "Every view ends in a dollar or a fan outcome."
- "AI-eval partner evaluations plus Dynatrace observability, now one company."

Don't say

- That we enforce agent identity or find Shadow IT
- An unverified kill switch or a private-beta feature
- Which LLMs Acme uses. Ask instead.

WHERE WE STAND

A former customer, a new buyer, a different Dynatrace

The 2017 loss was a commercial and coverage failure, not a capability verdict. Today's buyer is a different person with a different problem: **governing and proving AI**. The AE calls 2016/17 "a different Dynatrace at that time."

Relationship timeline

Oct–Nov 2014 · POC
Full-stack POC. Won as a **consolidation** of New Relic + Datadog (+ Pingdom).

Jan 2015 · New logo
36-month contract. TCV $126K / **ACV $42K** (24 host units) + $10K Quick Start.

Aug–Dec 2017 · Expansion lost
DEM POC **succeeded** but lost to a point RUM tool on price ($50K ACV). Add-on host units lost as well. They were over their HU limit.

Nov 18, 2017 · Cancellation
"Going another direction on price, and on coverage of the AWS services we use." The renewal closed lost on Dec 31. Internally it was called "a big surprise".

2017–2026 · Datadog era
Datadog is the incumbent. AE read (Jan 2026): *"doing alright with it."* A reseller "Competitive Displacement" opp has sat in Stage 1 since Mar 2025.

2025 · New AI leader
Marcus Hale joins to lead AI at Acme.

Mar 12, 2026 · Discovery
36-minute call that came from an AWS survey. AI opp created (Stage 2, AI Sales Play, Datadog incumbent).

Apr 8, 2026 · Written agenda
Hale declines a pre-call and emails 6 demo asks instead. Demo moved to **Apr 27**.

Why we lost, and why it's different now

2017 FAILURE	2026 POSTURE
Price vs. point tools (a point RUM tool)	DPS consumption. Land a scoped AI workload, not a host count.
"Not keeping up with AWS services"	Deep AWS integrations, Bedrock/Agentcore, the partner AI gateway work on AWS
Host-unit cap friction	Avoid caps in the land. Agree a ramp before overage hits.
Ops-only champion (Tech Ops)	AI leader with an executive-ROI mandate

Stakeholder map

PERSON	ROLE	STATUS
Marcus Hale	Director, AI · technical buyer	ENGAGED
Elena Ruiz	VP, Digital & Fan Experience	TARGET SPONSOR
Sam Okafor	Head of Event Technology	TARGET
Priya Nair	Sr Dir, Network / InfoSec	AGENT-IAM GATE
Greg Dorsey	Dir, IT Engineering (the reseller opp)	DATADOG OWNER?
Lena Brandt	Enterprise Technology Architect	UNENGAGED
Owen Pryce	Former champion (2014–17)	LIKELY LEFT

THE CUSTOMER WROTE THE AGENDA

Hale's six asks, and how we answer each one

Hale put his requirements in writing on Apr 8. That is our Apr 27 rubric, with one constraint: *"tie the telemetry to a business result, not just show AI activity for its own sake."*

HIS ASK (VERBATIM, ABRIDGED)	WHAT IT REALLY MEANS	DYNATRACE ANSWER	PROOF TO SHOW LIVE
1 "Performance, reliability and **cost** of AI-powered applications once they move into production"	Latency, error rate and token spend per app and feature	AI Observability: OpenTelemetry / OpenLLMetry GenAI spans, token & cost tracking, Dynatrace Intelligence anomaly detection	AI app health + $/feature view
2 "Visibility into AI agents and workflows that may span **multiple systems or models**… where diagnosing failures can otherwise become difficult"	Distributed tracing across agent → tool → model hops	End-to-end traces, Smartscape topology, causal root cause across services and LLM calls	Break a multi-agent flow live and find the root cause in minutes
3 "How **governance, observability and security** fit together as AI usage becomes more distributed"	Audit trail of agent actions, guardrails, data protection	Guardrail/eval signals on spans, PII/PHI redaction at the edge (OpenPipeline / Bindplane), AppSec. IAM enforcement via partners.	Governance view: who or what called what, with PII masked
4 "Examples where Dynatrace helped an organization move from **experimenting with AI to operating it reliably at scale**"	Social proof that matches his maturity curve	AI Observability customer references; Forrester TEI (466% ROI, <6-month payback)	2–3 reference stories, one in sports/media if available
5 "Connect the technical telemetry back to a **business outcome**"	An ROI narrative he can take to his executives	Business events + AI telemetry in Grail; executive dashboards	"AI Value Scorecard": spend, adoption and outcome KPI
6 "Seeing what **'good' looks like** once these capabilities are running in a real enterprise environment"	A target state and a maturity model	Reference architecture and an AI observability maturity ladder	A one-slide "good" end state with Acme placed on it

Additional asks raised verbally on the Mar 12 call

QA layer for vibe-coded software

He wants a gate that confirms code "won't allow anything malicious" and "will run properly" before production. (19:43)

A kill switch

"Does this include like a kill switch if need be?" The AE said "100 percent" but deferred the details to a technical resource. (22:24)

Identity & access management for agents

"We need just as much identity control on the agents as we have on the people." (21:40)

Agent productivity, not just health

"Is there anything that tells me not just whether an agent is healthy, but whether it's actually productive?" The AE answered "I don't know". (28:02)

WHAT WAS SAID

Explicit business problems, outcomes and technical needs

A direct read of the transcript, Hale's email and his public statements. Everything here is **stated by Acme**, with timestamps from the Mar 12 call.

BUSINESS PROBLEMS

> "I'm the only one here with AI in my title… it's a small team."
> Hale · 6:29

> "We're almost at like three people for every agent… and that flips next year, probably."
> Hale · 22:56

> "You just spent like eight grand in credits, so what did we get for it?… executives are expecting traditional metrics like ROI."
> Hale · 27:15

> "The Acme audience in particular is very sensitive about AI… they hold our feet to the fire."
> Hale · 3:04

- Rising autonomy brings "risk… challenges… higher levels of governance"
- Each department risks "reinventing the wheel" on AI workflows

DESIRED OUTCOMES

- **Measurable AI ROI** that executives accept, moving from "any use" to "high-value use"
- **Responsible AI** that protects fan trust and creative integrity
- **AI "starter license"**: upskilling plus a playbook per department
- **Scale output through AI**: a small team doing much more
- **Event ML**: pit-lane safety analytics and crowd flow at venues
- **Broadcast AI**: telemetry displays and racing lines on TV
- Agentic AI managed like staff

TECHNICAL SOLUTIONS SOUGHT

- Agent **health + productivity** monitoring
- **Identity & access** controls for agents
- Pre-production **QA / guardrail layer** for AI-generated code
- **Kill switch** for autonomous agents
- **LLM-agnostic** coverage ("doesn't matter which one we're using")
- Cross-system, multi-model **failure diagnosis**
- **Cost visibility** for AI in production
- Vendor must pass a single intake gate (SOC 2 required)

Three AI arenas at Acme, each with a different buyer

ARENA	EXAMPLES (STATED)	LIKELY OWNER	DYNATRACE RELEVANCE
Corporate / internal	AI starter license, department playbooks, internal comms, editorial, HR, social	Hale	PRIMARY · LAND AI value + governance
On-track / event	Pit-lane safety analytics, crowd flow at venues	Racing Ops / Event Technology	EXPAND real-time reliability
Broadcast / TV	Telemetry graphics, racing lines, GenAI lap summaries	Event Tech / Acme Broadcast	EXPAND live-data pipelines

WHAT WASN'T SAID, AND HOW WE LOSE

Reading the call backwards

Four lenses: **absence**, **pre-mortem**, **customer-side** and a **claim audit**.

Absence: signals that were missing

- **No infrastructure talk at all**: no cloud, datacenter, hypervisor or Kubernetes
- **Acme never mentioned Datadog.** The AE raised it. There's no stated observability pain.
- **No budget, timeline or decision process** discussed
- **No other stakeholders.** He'll only "consider" inviting others.
- **Wouldn't name his LLM stack** ("doesn't matter which one"). Public signals point to one major vendor.

Read: this is an **AI-governance and value buyer, not an ops buyer**.

Pre-mortem: how this deal dies

- **Generic platform tour.** That is exactly the "AI for its own sake" he warned against.
- **Over-promise caught** by his "AI intake lead" vendor vetting (SOC 2, single intake)
- **Single-threaded** on a director 9 months at Acme, with no ops budget, while IT defends Datadog
- **A replay of 2017 price shock**: we lost twice on price to cheaper point tools
- **Muddled AI-eval partner story**

Customer-side: what makes Hale the hero

- He's a public voice on AI (conference panels, an enterprise-AI summit, trade press). He needs **a story his executives believe**.
- "What did $8,000 of credits buy?" → give him a **repeatable ROI artifact**.
- His posture is governance first ("AI intake lead"). **Make Dynatrace the evidence layer** behind his intake gate.
- Small team → **automation over headcount**.

Strategic inversion: don't sell observability

The obvious play is "win back the APM seat from Datadog." Turned around: **sell an AI control and value system that happens to be built on observability**. Displacement follows once the platform is inside.

- Land: AI workloads (net new, no incumbent fight)
- Expand: logs and cost (Bindplane) → full stack via the reseller opp

Claim audit: commitments we must back up before Apr 27

CLAIM MADE TO ACME	WHERE	STATUS	ACTION
The partner AI gateway has a kill switch ("100 percent certain")	Call, 22:32	UNVERIFIED	Confirm with product. Show the real stop or rollback path (AutomationEngine, feature flags).
Dynatrace plays in agent identity & access ("absolutely")	Call, 23:47 · Deck	PARTIAL	Reframe: **observe and audit** agent identity and actions. Enforcement sits with IAM partners.
Agent productivity / ROI metrics	Deck	PARTIAL	Build it with business events plus GenAI spans.
"Works across *your* models"	Deck	MISATTRIBUTED	That was **Dynatrace's** list. Correct to "LLM-agnostic".

KNOWN VS. UNKNOWN

Acme's technology, AI and datacenter / hypervisor stack

The AI layer is reasonably clear. The **datacenter and hypervisor layer has no evidence in any source**. That makes it the top discovery priority for Apr 27.

LAYER	WHAT WE KNOW	CONFIDENCE	SOURCE
LLM providers	Leans toward one major LLM vendor (public talks). An enterprise chat assistant is rolled out to staff. Requires LLM-agnostic tooling.	MEDIUM	LinkedIn, press, call
Agents	About 1 agent per 3 employees, expected to exceed headcount in 2027. Agentic autonomy is rising.	HIGH	Hale, call
Software delivery	"Vibe coding is big." AI-generated code is reaching production.	HIGH	Hale, call
AI governance	Single "front door" intake ("AI intake lead"), SOC 2 required of vendors, agents managed "like employees"	HIGH	Trade press
Data platform	Lakehouse signals: those skills show up in Acme's data job posts.	LOW–MED	Job posts
Cloud	AWS is the primary cloud (public). The 2017 churn cited "AWS services". This engagement came from an AWS survey.	MEDIUM	Salesforce, press, call
Observability	Datadog incumbent again since 2017. A point RUM tool for RUM/crash (2017). New Relic and Pingdom historically.	MEDIUM	Salesforce, Slack, call
Real-time / broadcast	Race telemetry, GenAI lap summaries, live streaming	MEDIUM	acmeracing.example, press, LinkedIn
Endpoint / network	Several thousand managed devices (IT engineering). Dedicated network and InfoSec leadership.	LOW	LinkedIn
Datacenter / hypervisor	**No evidence found.** INFERRED: event venues (timing, ticketing, broadcast compounds) and league HQ likely keep on-prem or edge compute. VMware vs. Nutanix vs. Hyper-V is unknown.	UNKNOWN	Gap: validate

Discovery questions: AI stack

- Where do AI workloads run: OpenAI API, Azure OpenAI, Bedrock / Agentcore, Databricks?
- What builds the agents: Custom GPTs, Agents SDK, LangGraph, low-code?
- Who owns the AI spend number today, and how is it tracked?
- Which AI workloads face fans, and which are internal?

Discovery questions: datacenter / hypervisor

- What stays on-prem at tracks and league HQ? Which hypervisor runs it (VMware / Nutanix / Hyper-V)?
- Is any race-control or broadcast compute at the edge, and on what platform?
- What does Datadog cover today, what doesn't it cover, and what does it cost?

LINKEDIN + PUBLIC WEB · LAST 60 DAYS, PLUS KEY CONTEXT

What Acme is saying about AI in public

The public story matches the call: **governance first, ROI over hype, trust over access**. Hale is building a profile as a sports-AI voice. Fan-facing AI mostly comes from partners on AWS, and the league has been burned by visible AI before.

Last 60 days

DATE	SIGNAL
~Mar 30	**LinkedIn · Hale**: sports-tech conference panel. Takeaway: "trust grows when people can see how AI helps with their own work."
Mar 18	**Trade press**: Hale: "The question is whether they'll use it well." He warns of **Shadow IT**, staff using unapproved AI tools.
~Mar	**LinkedIn · Hale**: enterprise-AI summit talk on moving from **"any AI use" to "high-value use."**
Mar 5	**Industry press**: the Digital VP says a fan-facing AI experiment "missed the mark." AI stays behind the scenes.
~Feb	**Trade press (feature)**: Hale as the league's **"AI intake lead"**. Single "front door" intake, **SOC 2 required**, agentic AI treated like staff, not "AI for its own sake."

Earlier context

Jan 2026	Enterprise chat-assistant rollout to staff. Analytics lead: fan-feedback analysis dropped from hours to minutes.
2025	Hale joins to lead AI at Acme

> "The goal is to give every person more reach, not fewer colleagues."
>
> Marcus Hale · trade press

> "People adopt what they trust, not what they're handed."
>
> Marcus Hale · conference panel

Fan-sentiment landmines

- A fan-facing AI content experiment drew public criticism

Implication: fan-facing AI must be invisible and reliable. Pitch confidence, not novelty.

Infrastructure signals (public)

- **AWS** is the primary cloud
- Race-data platform modernized to cloud, with edge compute at events

What this means for Dynatrace

→ **Be the intake lead's instrument panel**: which AI is used, by whom, with what data, at what cost.

→ **AWS co-sell**: Marketplace purchase can draw down committed spend.

→ **The chat assistant is the starting estate.** Agentic workloads are where observability becomes mandatory.

→ **Race-day expansion**: edge and cloud race-data telemetry. A competition-integrity play for later.

MONDAY, APRIL 27 · 10–11 AM PT

The demo: "From AI activity to AI value"

One narrative, run in his order, on a **Acme-flavored AI workload**: an illustrative "Race Weekend Fan Assistant" agent that calls an LLM plus tools (schedule, parking and ingress, tickets), instrumented with OpenTelemetry. Every segment ends with the **business outcome**, not the feature.

TIME	SEGMENT	WHAT HE SEES	ASK
0–5	**His agenda, on screen**	Replay his 6 asks and agree what "success today" looks like	—
5–15	**The AI value scorecard**	Exec dashboard: AI spend, tokens and cost per feature, adoption, and one outcome KPI (e.g. fan queries resolved without an agent)	1 5
15–28	**Anatomy of an agent failure**	Multi-agent, multi-model trace. Inject a fault (model timeout or tool error) and let Dynatrace Intelligence point to the root cause.	2
28–38	**Governed by design**	Guardrail and eval signals, PII masked at the edge, an audit of which agent called what. Show the **verified** stop / rollback path.	3
38–46	**Safe speed for vibe-coded software**	AI-generated change → production-aware review → automated guardrail (AI SRE / the partner AI gateway, once verified)	Verbal
46–54	**What "good" looks like**	Maturity ladder with Acme placed on it, plus 2–3 experiment-to-scale customer stories	4 6
54–60	**The proposal**	The 30-day AI Value & Governance POV (next page) and named next-step owners	—

Do

- Lead with his words and his order
- Tie every view to a dollar or a fan outcome
- Leave him a screenshot-ready ROI artifact
- Bring the SOC 2 / trust package

Don't

- Pitch platform consolidation or bash Datadog
- Claim agent IAM enforcement or an unverified kill switch
- Name Acme's LLMs for them
- Run a feature tour

Who

- **Alex Moreno**: AI narrative and live demo
- **Jordan Reyes**: business framing, close
- **Casey Lin**: warm exec intro
- **Riley Park** (BDR): outreach
- Principal SE on standby for the POV

PROPOSED PROOF OF VALUE

30-day AI Value & Governance POV

One real AI workload, five measurable criteria and an exec readout **Hale presents himself**. Net-new scope, so there's no incumbent fight. It shows the value that justifies the later platform conversation.

Scope

- **One AI workload** in production or late pre-prod. Candidates: the internal enablement assistant, the editorial/social content workflow, or a fan-facing assistant.
- OpenTelemetry / OpenLLMetry instrumentation across at least 2 models or tools
- Business events for one outcome KPI
- Optional: edge PII redaction in the telemetry pipeline

Success criteria (agreed in week 0)

#	CRITERION	MAPS TO
1	100% of LLM and agent calls traced end-to-end	Ask 2
2	Daily cost per workflow / feature on an exec dashboard	Ask 1, 5
3	An agent failure diagnosed in minutes, not hours (vs. baseline)	Ask 2
4	PII masking verified and policy violations alerted	Ask 3
5	One business KPI tied to AI telemetry	Ask 5, ROI

Plan

Week 0 · Align
Success criteria, data-handling review, trust package through his intake gate

Week 1 · Instrument
OTel on the chosen workload. First traces and cost data on day one.

Weeks 2–3 · Prove
Value scorecard, governance view, a fault drill

Week 4 · Readout
Hale presents results to his executives. We present the land proposal.

Commercial path

- **Land**: DPS, scoped to AI workloads. No host caps.
- **Fast-land option**: a 2-week guided value sprint, or a direct land with a success plan (the AE wants "ideally no proof")
- **Expand**: logs and cost (Bindplane), then full stack through the reseller "Competitive Displacement" opp
- Engage the **AI-eval partner** counterpart per the account-mapping guidance

BEFORE APRIL 27

Next steps, owners and leadership asks

ACTION	OWNER	BY
Verify the kill switch, the partner AI gateway and agent identity capabilities with product. Fix the claims accordingly.	Alex Moreno	Apr 17
Complete AI-eval partner account mapping and pick AI target accounts	Jordan Reyes / Casey Lin	Apr 17
Correct the "What We Heard" slide (models are Dynatrace's, not Acme's). Add the "LLM-agnostic" framing.	Jordan Reyes	Apr 21
Build the league-themed AI demo workload (fan-assistant agent, value scorecard, governance view)	Alex Moreno	Apr 24
Assemble the trust package (SOC 2, data handling, AI governance) for his vendor intake	Jordan Reyes	Apr 24
Pull 2–3 experiment-to-scale AI customer references (sports/media preferred) and the Forrester TEI	Jordan / Alex	Apr 24
Ask Hale to invite AI-first peers. Run targeted outreach to the Digital VP, Event Technology and InfoSec.	Jordan / Riley	Apr 24
Align with the reseller on the "Competitive Displacement" opp and the proposed on-site visit	Jordan Reyes	May 1
Confirm the AWS-survey thank-you (a branded gift) reached Hale	Riley Park	Now

Ask 1 · Executive sponsor

A Dynatrace executive paired with a Acme executive (Digital VP level) after the demo. The AE wants an exec-to-exec thread out of Apr 27.

Ask 2 · Principal SE coverage

An AI-observability principal on standby to run the POV if it's accepted. It's a marquee sports logo.

Ask 3 · Commercial flexibility

An AI-scoped DPS land with no caps. We lost twice on price in 2017. Don't repeat it.

ONE-LINE SUMMARY FOR LEADERSHIP

Acme's AI leader handed us the scorecard. If we answer it honestly and tie it to dollars, we can land AI observability inside a former customer and earn the right to win the platform back.

Sources

Mar 12, 2026 discovery call transcript (Gong) and SE leadership call summary · Hale email, Apr 8, 2026 · Salesforce (Acme account, opps, contracts, POCs; queried Apr 8) · Microsoft 365 mail and calendar · Slack DMs (Jan–Apr 2026) · Highspot · LinkedIn (Hale and Acme leaders, posts since Feb 7, 2026) · motorsport trade press · acmeracing.example · public web (see External Intelligence). Items marked INFERRED are analyst judgment, not stated fact.